

## Table of Contents



### ▼ 1 Initialization cell -- run this first!

```
In [3]: 1 using Colors, Interact, Plots, JLD, LinearAlgebra, Arpack
2 gr(
3     label="",
4     markerstrokewidth=0.3,
5     markerstrokecolor="white"
6 )
```

Run

Out[3]: Plots.GRBackend()

### ▼ 2 Low rank matrices

A set of vectors is said to be *linearly dependent* if one of the vectors in the set can be expressed as a linear combination of the others. If no vector in the set can be written as a linear combination of the others, then the vectors are said to be *linearly independent*.

Concretely, suppose  $a_1, \dots, a_n$  are a set of vectors. Then they are considered to be linearly dependent if there exist scalars  $\alpha_1, \dots, \alpha_n$  (not all zero) such that

$$\alpha_1 a_1 + \dots + \alpha_n a_n = 0. \quad (1)$$

If no such scalars exist, then the vectors are linearly independent. In other words, if the vectors are linearly independent, then the only way to satisfy the equation above is to set  $\alpha_1 = \dots, \alpha_n = 0$ .

The vectors  $a_1 = [1, 0]^T$  and  $a_2 = [2, 0]^T$  are linearly dependent because  
 $a_2 = 2a_1 \Rightarrow a_2 - 2a_1 = 0$ .



True

False

Submit





The vectors  $a_1 = [1, 0]^T$  and  $a_2 = [0, 1]^T$  are linearly dependent



- True  
 False



The **column rank** of an  $m \times n$  matrix  $A$  is the number of linearly independent columns.

The **row rank** of an  $m \times n$  matrix  $A$  is the number of linearly independent rows.

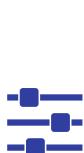
A remarkable fact of linear algebra is that the number of linearly independent rows of a matrix equals the number of linearly independent columns:

$$r = \text{rank}(A) = \text{rank}(A^T) \quad (1)$$

Since the number of linear independent columns (or rows) cannot exceed the number of columns (or, resp., rows) in a matrix, this implies that

$$r \leq \min(m, n). \quad (2)$$

$A$  is a  $10 \times 5$  matrix. It can have rank (select **all** that apply)



- 10  
 7  
 5  
 2  
 1



Consider the matrix

$$A = xy^T,$$

where  $x$  is an  $m$ -dimensional vector, and

$$y = [y_1 \quad \dots \quad y_n]^T.$$

Consequently, the matrix  $A$  can be written in column partition format as

$$A = [y_1 x \quad \dots \quad y_n x].$$

We would like to determine the rank of  $A$ .

Assume  $x, y \neq 0$ , and there is an index  $i$  such that  $y_i \neq 0$ . The rank of  $A = xy^T$  is 1 because (select **all** that apply)

- $1 \cdot A[:, j] - (y_j/y_i) \cdot A[:, i] = 0$  for every  $j = 1, \dots, n$  so the columns of  $A$  are linearly dependent





- Every column of  $A$  is a scalar multiple of its  $i$ -th column  $y_i x$ . Thus the  $i$ -th column is the only independent column -- the rest depend on it.
- $A$  as written in column partition format "looks" like it has just one row

Submit

Now suppose that  $x_1$  and  $x_2$  are linearly independent vectors and

$$A = x_1 y_1^T + x_2 y_2^T.$$

What is the rank of  $A$ ? Let us begin by examining the structure of  $A$  -- as earlier, it can be expressed as

$$A = [y_{11}x_1 \quad \dots \quad y_{1n}x_1] + [y_{21}x_2 \quad \dots \quad y_{2n}x_2],$$

or equivalently

$$A = [y_{11}x_1 + y_{21}x_2 \quad \dots \quad y_{1n}x_1 + y_{2n}x_2].$$

We can now observe that every column of  $A$  is a linear combination of two linearly independent vectors. Thus the rank of  $A$  is at most two. Note now that

$$A^T = y_1 x_1^T + y_2 x_2^T,$$

and so we can see that the rank of  $A$  also depends whether  $y_1$  and  $y_2$  are linearly independent. If they are linearly dependent (and not equal to zero), then the rank of  $A^T$  (and hence  $A$ ) will be 1!

Being able to spot low-rank matrices via the dependency structure in the rows or columns is a valuable skill in data science, machine learning and computational mathematics.

Suppose

$$x_1 = [1, 0]^T, \quad x_2 = [0, 1]^T, \quad y_1 = [1, 0]^T, \quad y_2 = [1, 1]^T.$$

Then the rank of  $x_1 y_1^T + x_2 y_2^T$  is

- 2, because  $x_1$  and  $x_2$  are linearly independent, and so are  $y_1$  and  $y_2$ .
- 1, because  $y_1$  and  $y_2$  are linearly dependent.

Submit

Suppose

$$x_1 = [1, 0]^T, \quad x_2 = [2, 0]^T, \quad y_1 = [1, 0]^T, \quad y_2 = [1, 1]^T.$$

Then the rank of  $x_1 y_1^T + x_2 y_2^T$  is

- 1, because  $x_1$  and  $x_2$  are linearly dependent.
- 1, because  $y_1$  and  $y_2$  are linearly dependent.
- 2, because  $x_1$  and  $x_2$  are linearly independent, and so are  $y_1$  and  $y_2$ .

Submit

Suppose

$$x_1 = [1, 0]^T, \quad x_2 = [1, 1]^T, \quad x_3 = [0, 1]^T, \quad y_1 = [1, 0]^T, \quad y_2 = [1, 1]^T, \quad y_3 = [2^{-1/2}, 3]$$

Then the rank of  $x_1y_1^T + x_2y_2^T + x_3y_3^T$  is (select all that apply)



- at most 2, because it is a  $2 \times 2$  matrix.
- 2, because  $x_3$  is linearly dependent on  $x_1$  and  $x_2$ , but  $x_1$  and  $x_2$  are linearly independent.
- 2, because  $y_3$  is linearly dependent on  $y_1$  and  $y_2$ , but  $y_1$  and  $y_2$  are linearly independent.
- 1, because  $x_3$  is linearly dependent on  $x_1$  and  $x_2$ .
- 1, because  $y_3$  is linearly dependent on  $y_1$  and  $y_2$ .

Submit

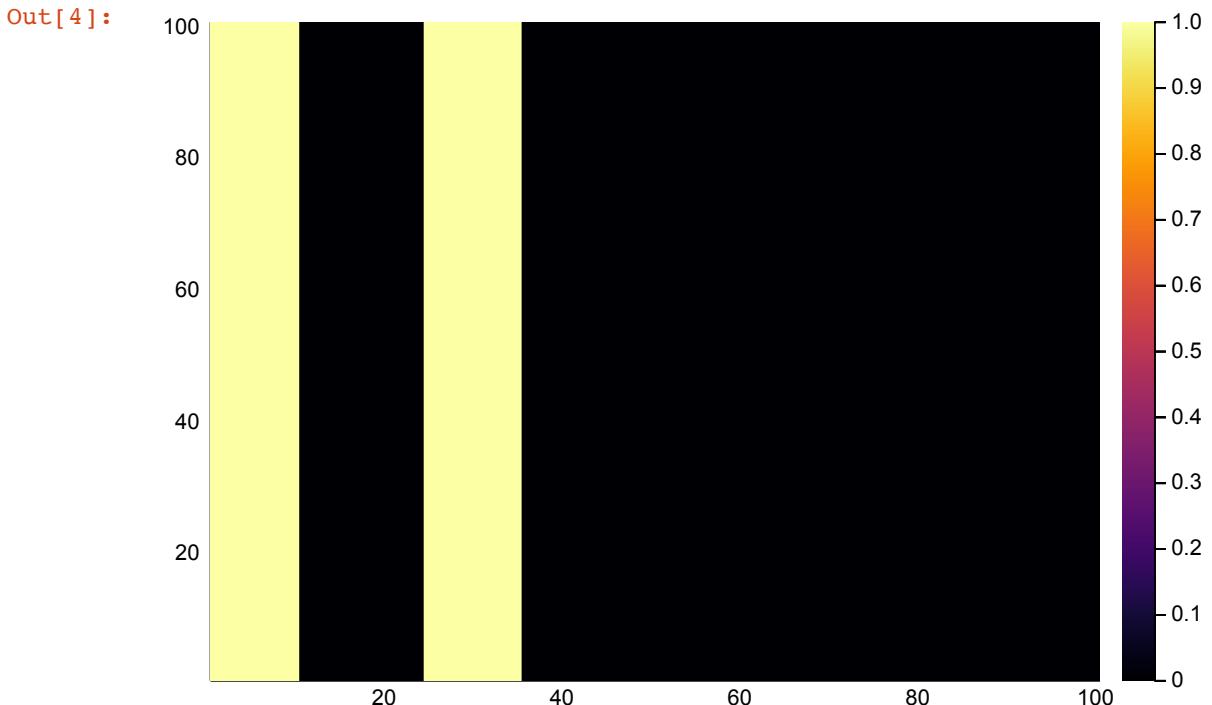
## ▼ 2.1 Spotting low-rank matrices

The ability to spot low-rank matrices "in the wild" is a useful skill in computational data science and applied mathematics. Let's exercise that muscle with some warm-ups :-)

**Tip:** For the questions that follow, express the matrix as  $xy^T$  or  $x_1y_1^T + x_2y_2^T + x_3y_3^T$ , then examine the dependency structure of the  $x$ 's and  $y$ 's to deduce the rank of the matrix.

```
In [4]: 1 A = [ (1 <= j <= 10) || (25 <= j <= 35) ? 1 : 0
          2     for i in 1:100, j in 1:100 ]
          3 heatmap(A)
```

Run



The rank of  $A$  is



- 100, because there are 100 rows and columns in the matrix.



- 2, because there are two non-zero columns.
- 1, because the same column is repeated in the other non zero columns.

[Submit](#)



A can be expressed as the outer product  $x * y'$ , where

- $x = \text{ones}(100)$  and  $y = [\text{ones}(10); \text{zeros}(15); \text{ones}(10); \text{zeros}(65)]$ .
- $x = \text{ones}(100)$  and  $y = [\text{ones}(10); \text{zeros}(14); \text{ones}(11); \text{zeros}(65)]$ .

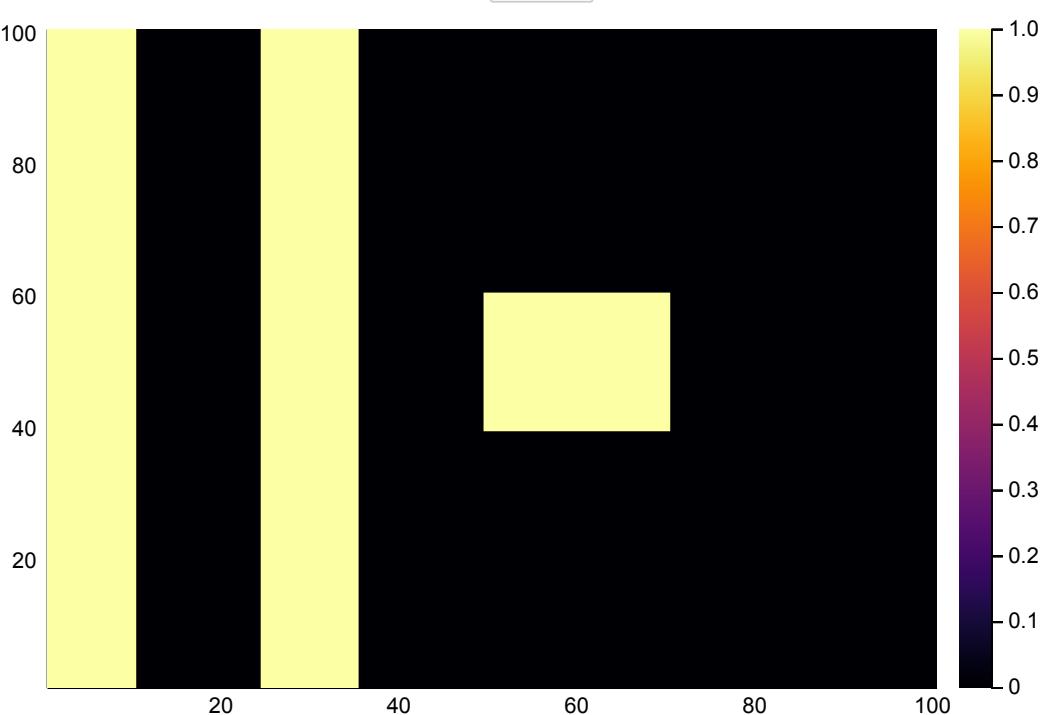
[Submit](#)

In [5]:

```
1 A = [ (1 <= j <= 10) || (25 <= j <= 35) ||  
2     (40 <= i <= 60 && 50 <= j <= 70) ? 1 : 0  
3     for i in 1:100, j in 1:100 ]  
4 heatmap(A)
```

[Run](#)

Out[5]:



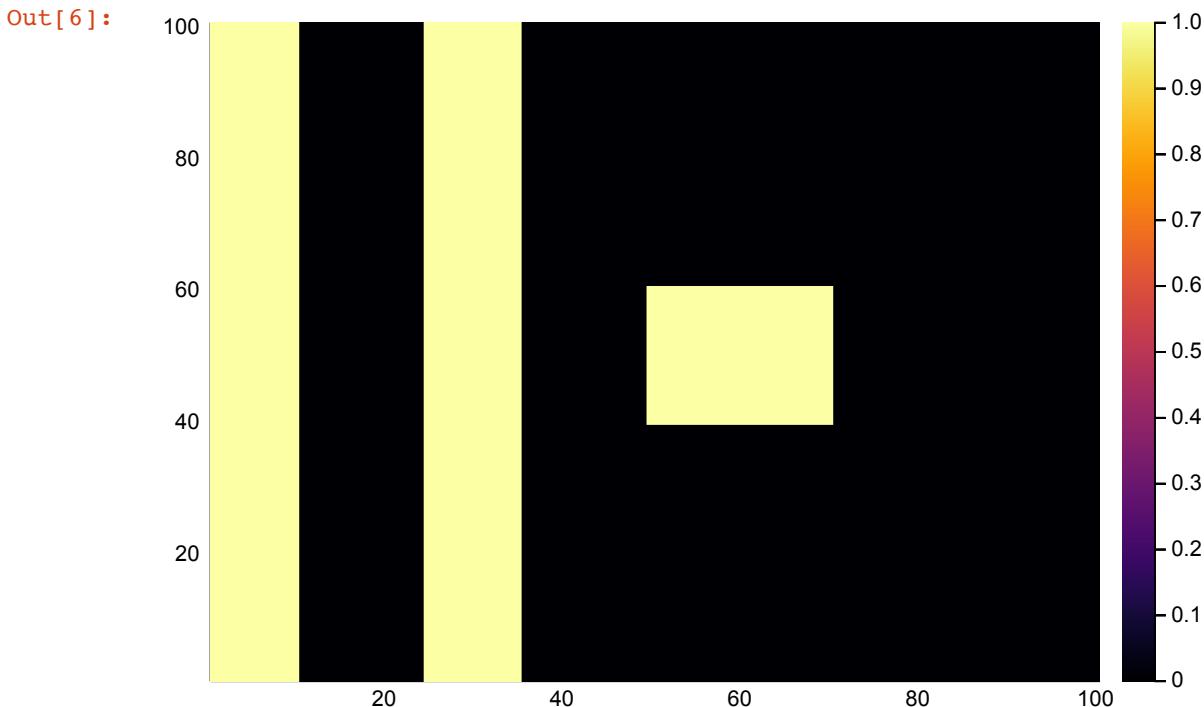
A can be expressed as  $x_1 * y_1' + x_2 * y_2'$  where (select all that apply)

- $x_1 = \text{ones}(100)$  and  $y_1 = [(25 <= j <= 35) || (1 <= j <= 10) ? 1 : 0 \text{ for } j \text{ in } 1:100]$ .
- $x_2 = \text{ones}(100)$  and  $y_2 = [50 <= j <= 70 ? 1 : 0 \text{ for } j \text{ in } 1:100]$ .
- $x_1 = \text{ones}(100)$  and  $y_1 = [50 <= j <= 70 ? 1 : 0 \text{ for } j \text{ in } 1:100]$ .
- $x_2 = [(40 <= i <= 60) ? 1 : 0 \text{ for } i \text{ in } 1:100]$  and  $y_2 = [50 <= j <= 70 ? 1 : 0 \text{ for } j \text{ in } 1:100]$ .

[Submit](#)

Indeed! Let us check it just to verify.

```
In [6]: 1 x1 = ones(100)
2 y1 = [(25 <= j <= 35) || (1 <= j <= 10) ? 1 : 0 for j in 1:100]
3 x2 = [(40 <= i <= 60) ? 1 : 0 for i in 1:100]
4 y2 = [50 <= j <= 70 ? 1 : 0 for j in 1:100]
5 Aalt = x1 * y1' + x2 * y2'
6 heatmap(Aalt)
```



We can compute the Frobenius norm of the difference between the two matrices using the command `vecnorm` as in the next cell.

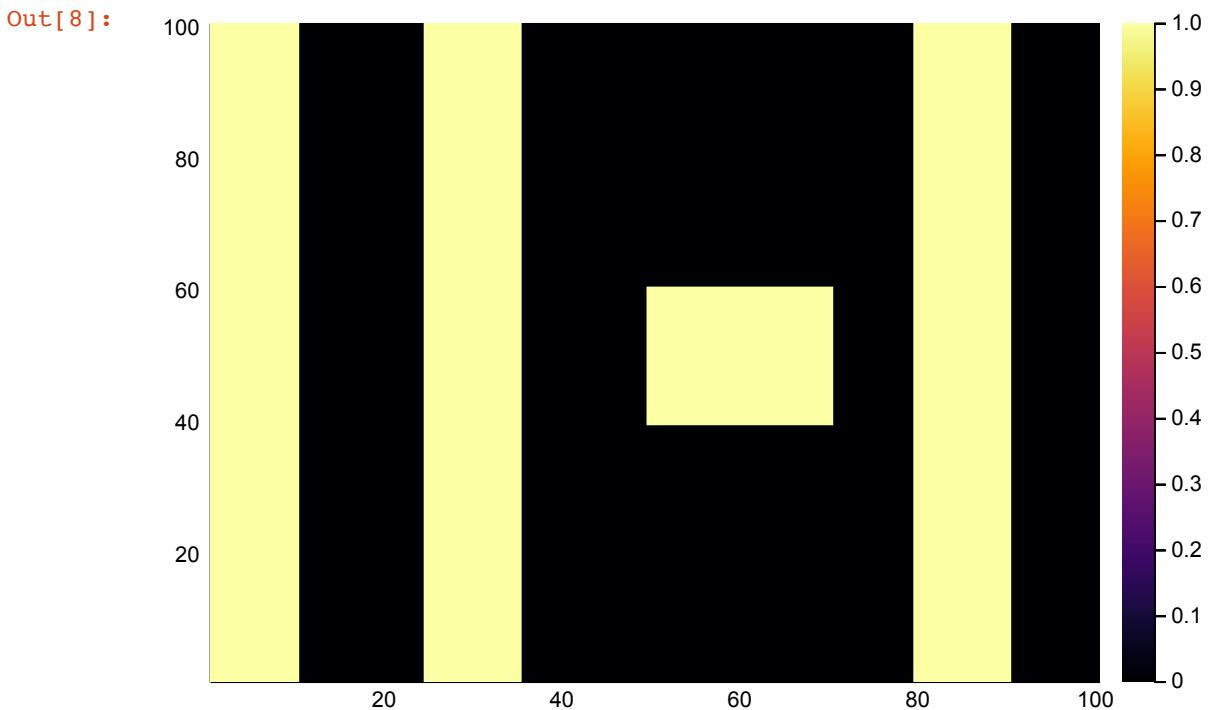
```
In [7]: 1 @show norm(A - Aalt);
```

```
norm(A - Aalt) = 0.0
```

A few more exercises and we will be ready for the main point of this section -- which is to use the SVD extract low-rank matrices!

```
In [8]: 1 A = [(1 <= j <= 10) || (25 <= j <= 35) ||
2           (40 <= i <= 60 && 50 <= j <= 70) ||
3           (80 <= j <= 90) ? 1 : 0
4           for i in 1:100, j in 1:100]
5 heatmap(A)
```

Run



What is the rank of  $A$ ?

- 1
- 2
- 3
- 4

Submit

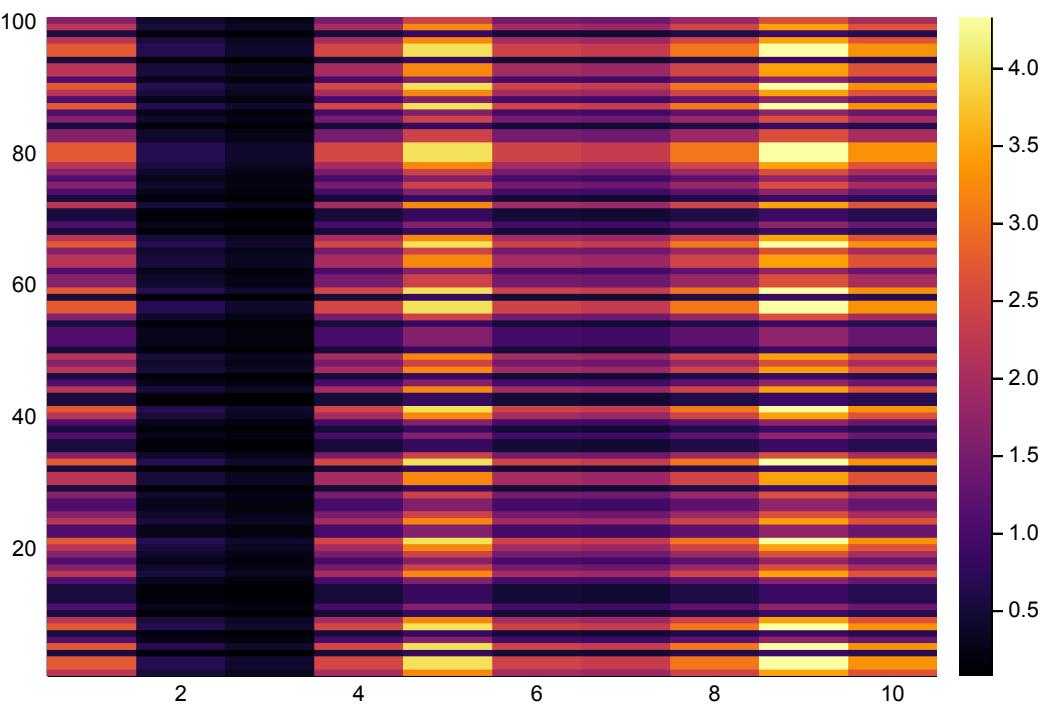
Indeed! The matrix  $A$  can be written as  $x_1 y_1^T + x_2 y_2^T$ . The  $x_1$  and  $x_2$  vectors are same as earlier. Only the  $y_1$  entry changes -- we add extra ones in the indices of  $y_1$  corresponding to the columns where the new "stripe" appears.

Now consider the matrix in the following cell.

```
In [9]: 1 x = rand(1:5, 100)
2 y = rand(1, 100)
3 A = [y[j] * x[i] for i in 1:100, j in 1:10]
4 heatmap(A)
```

**Run**

Out[9]:



What is the rank of  $A$ ?



- 100, because there are 100 rows.
- 10, because there are 10 columns.
- 1, because the matrix can be written as  $xy^T$ , where elements of  $y$  are random numbers given by `rand`.

**Submit**

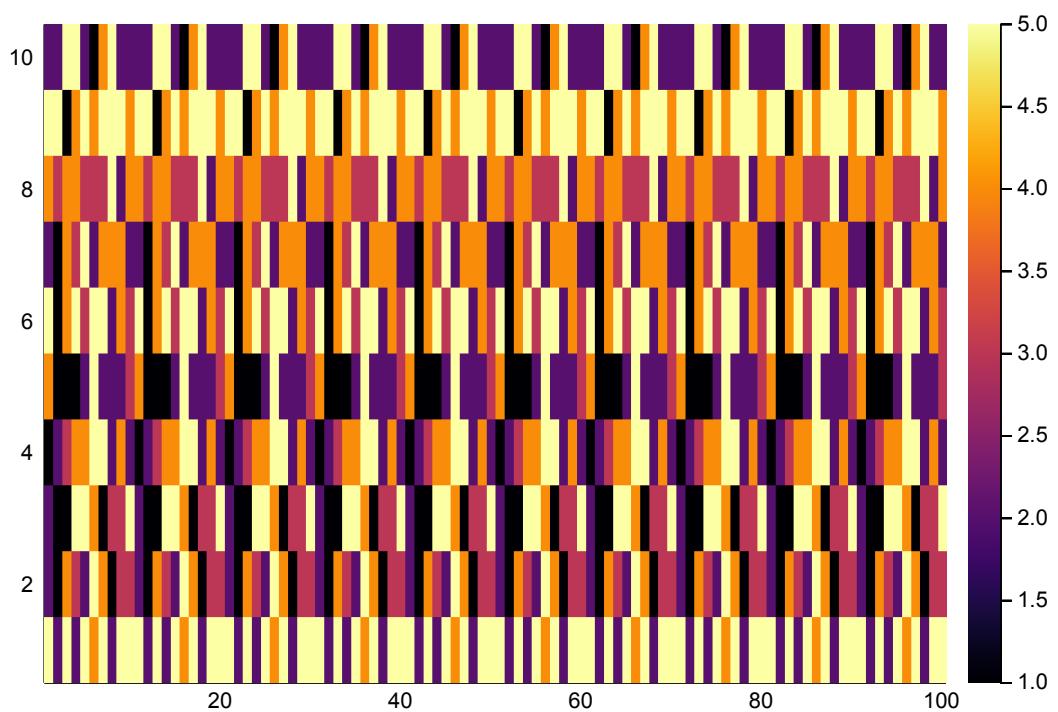
## 2.2 Spotting an *induced* rank-one matrix

Consider the matrix  $B$  in the following cell. It is obtained by concatenating copies of the matrix  $X$  alongside each other.

```
In [10]: 1 X = rand(1:5, 10, 10)
          2 B = repeat(X; outer=(1, 10))
          3 heatmap(B)
```

Run

Out[10]:

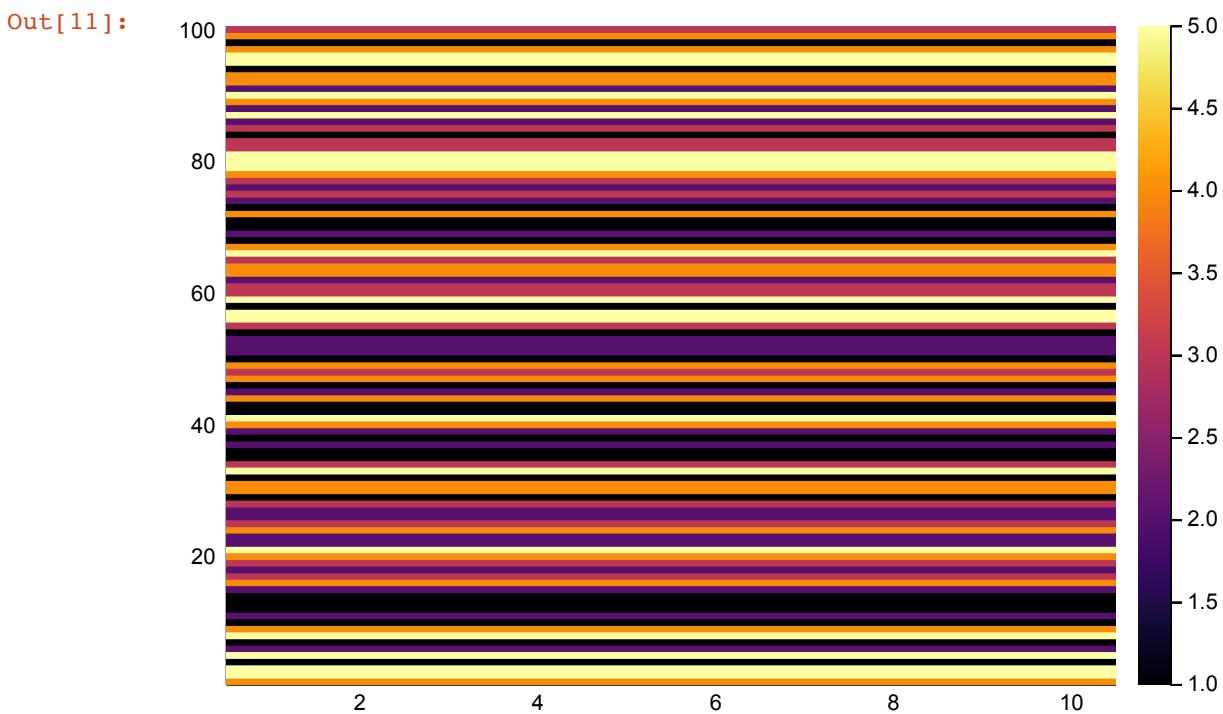


The columns of  $X$  are drawn at random, so generically (with very high probability) the rank of  $X$  will equal the number of columns. Thus, the matrix  $B$  will have as many linearly independent columns as the number of linearly independent columns of  $X$ .

Now consider the matrix  $A$  formed in the next cell.

```
In [11]: 1 A = repeat(x[:, :], outer=(1, 10))
2 heatmap(A)
```

Run



$A$  and  $B$  have the same elements of  $A$ , but in a different order. The rank of  $A$  is



- equal to the rank of  $B$ , regardless of the rank of  $B$ .
- 1, even when the rank of  $B$  is higher.

Submit

In this case, even though the matrix  $B$  did not have rank one, we **induced a low-rank matrix  $A$**  by reshaping the elements of  $B$  in an appropriate manner. This will be important for understanding why the SVD application we describe next works the way it does.



### 3 Application: Background subtraction using the SVD

SVD decomposes a matrix into a sum of rank-one outer product matrices. This powerful representation enables a seemingly magical feat when applied to stable video footage.

We will illustrate with a surveillance video, currently stored as a `.jld` file. The command `JLD.load` produces a Dictionary of the variables in the file.

In [12]: 1 `load("lobby.jld")`

 Run

Out[12]: Dict{String,Any} with 1 entry:  
`"movie_cube" => Float16[0.3452 0.341 ... 0.2627 0.2588; 0.447 0.443 ... 0.3254 0...]`

The command `load` is used to open a `.jld` formatted file. We can access a specific variable, `movie_cube` in this case, as follows:

In [13]: 1 `# TODO: Fill in the ?? with the appropriate variable name from`  
 2 `# the Dictionary`  
 3 `file = load("lobby.jld")`  
 4 `varname = "movie_cube"`  
 5  
 6 `movie_cube = file[varname]`  
 7  
 8 `# data are compressed using 16-bit containers;`  
 9 `# we prefer 64-bit for computation`  
 10 `movie_cube = Float64.(movie_cube)`  
 11  
 12 `# alternate way to load variable:`  
 13 `# movie_cube = load("lobby.jld", varname)`  
 14  
 15 `@show size(movie_cube)`  
 16 `@show typeof(movie_cube);`

 Run

`size(movie_cube) = (128, 160, 650)`  
`typeof(movie_cube) = Array{Float64,3}`

The `movie_cube` variable is of type



- `Array{Float64,1}` -- i.e. `Vector{Float64}`
- `Array{Float64,2}` -- i.e. `Matrix{Float64}`
- `Array{Float64,3}` -- it's a "data cube"

 Submit

The `movie_cube` datacube is composed of 650 frames of a video, where each frame is a 128-by-160 image. We will scale the elements of the datacube so that they lie between 0 and 1, then display the first frame of this video.

```
In [14]: 1 """
2     B = rescale_zero_one(A)
3     Return a shifted and scaled version of input array `A` where all
4     elements are in the range [0, 1].
5 """
6 function rescale_zero_one(A::Array)
7     B = float(A)
8     B .-= minimum(B)
9     B ./= maximum(B)
10    return B
11 end
12
13 movie_cube = rescale_zero_one(movie_cube)
14
15 # display first frame
16 ▾ Gray.(movie_cube[:, :, 1])
```

Run

Out[14]:



Use the cell below to explore different frames of the video.

```
In [15]: 1 function plotframe(X; kw_args...)
2     return heatmap(
3         X;
4         color=:grays,
5         yflip=true,
6         aspect_ratio=:equal,
7         border=false,
8         ticks=[],
9         axis=false,
10        colorbar=false,
11        kw_args...,
12    )
13 end
```

▶ Run

Out[15]: plotframe (generic function with 1 method)

```
In [16]: 1 nframes = size(movie_cube, 3)
2 @manipulate for frame = 100
3     plotframe(
4         movie_cube[:, :, frame],
5         title="Frame $frame of $(nframes)",
6         clim=(0, 1)
7     )
8 end
9
10 # Enter values in the box below to explore the video
```

▶ Run

Out[16]: frame

Frame 100 of 650



Let us now play the matrix as a movie:

```
In [17]: 1 include("./webplayer.jl")
2 playvideo(
3     [movie_cube], ["Original movie"];
4     frames_per_second=120
5 )
6 # move the slider to scan through the movie
```

Run

Out[17]: 325



What happens in the movie? (select **all** that apply)

- There are people walking around.
- The lights go on and off.
- The camera angle is fixed.
- Around frame 182, four lights are off and just one remains on.
- Around frame 450, the four lights that were off come back on.
- All the lights stay on throughout the scene.
- The middle light in the background stays on through the scene.

Submit

## ▼ 3.1 Decomposing the movie matrix with the SVD

We will now use the SVD to decompose the movie. The movie is stored as an  $m \times n \times t$  3-D array, where

- $m$  is the number of rows of pixels in each frame,
- $n$  is the number of columns of pixels in each frame, and
- $t$  is the number of frames in the movie

To decompose using the SVD, we first reshape the movie array into an  $mn \times t$  **matrix  $A$** . Each column of the matrix corresponds to a frame (and contains  $mn$  pixel values).

We then compute the rank- $k$  truncated singular value decomposition of  $A$ . In other words, if  $A$  has rank  $r$  and has SVD

$$A = \sum_{i=1}^r \sigma_i u_i v_i^T,$$

then  $A_k$ , the rank- $k \leq r$  truncated SVD of  $A$ , is

$$A_k = \sum_{i=1}^k \sigma_i u_i v_i^T,$$

and  $(A - A_k)$  is the residual matrix. We will compute the residual  $(A - A_k)$  directly, but it is good to observe that it has SVD

$$A - A_k = \sum_{i=1}^r \sigma_i u_i v_i^T - \sum_{i=1}^k \sigma_i u_i v_i^T = \sum_{i=k+1}^r \sigma_i u_i v_i^T.$$

In the cell below, write a function `svdize_cube` that takes as input the  $m \times n \times t$  movie, and a rank  $k$ . The function should

1. reshape the movie into an  $mn \times t$  matrix
2. compute the truncated SVD and residual
3. reshape both truncated SVD and residual into  $m \times n \times t$  arrays
4. return the reshaped truncated SVD, reshaped residuals, and singular values

**Hint 1:** Recall that the above sum of outer products is equivalent to

$$A_k = U_k \Sigma_k V_k^T$$

where  $U_k = [u_1 \cdots u_k]$  is a matrix of the first  $k$  left singular vectors,  $V = [v_1 \cdots v_k]$  is a matrix of the first  $k$  right singular vectors, and  $\Sigma = \text{diagm}(\sigma_1, \dots, \sigma_k)$  is a diagonal matrix of the first  $k$  singular values.

**Hint 2:** The matrices  $U_k$ ,  $\Sigma_k$  and  $V_k$  can be obtained using `svds` as

```
U, s, V = svds(A; nsv=k)[1]
```

```
In [18]: 1  using LinearAlgebra
2  using Arpack: svds
3
4  """
5      cube_k, cube_residual, sk = svdize_cube(cube, k)
6
7  Inputs:
8  * `cube` is an m x n x l array
9  * `k` is an integer
10
11 Outputs:
12 * `cube_k` is an m x n x l 3-D array
13 * `cube_residual` is an m x n x l 3-D array
14 * `sk` is a length k vector
15
16 Given a 3-D array `cube` and a rank `k`, return a (rank `k`)
17 truncated SVD version of the cube as `movie_cube_k`, the residual
18 as `cube_residual`, the right singular vectors as `V_k`,
19 and the singular values as `sk`.
20
21 *Hint: svds has a keyword argument `nsv` for the number of singular
22 values. To specify this argument, you use the keyword:
23 `svds(..., nsv=...)`*
24 """
25 function svdize_cube(cube, k)
26     # Reshape 3-D array into matrix where each column is a frame
27     m, n, l = size(cube) # cube is an m x n x l 3-D array
28     matrix = reshape(cube, m*n, l) # Convert to mn x l matrix
29
30     # Compute rank k truncated SVD and residual
31     UsV = svds(matrix; nsv=k)[1]
32     Uk = UsV.U
33     sk = UsV.S
34     V_k = UsV.Vt
35
36     matrix_k = Uk * Diagonal(sk) * V_k
37     matrix_residual = matrix - matrix_k
38
39     # Reshape back into 3-D m x n x numFrames arrays
40     cube_k = reshape(matrix_k, m, n, l)
41     cube_residual = reshape(matrix_residual, m, n, l)
42
43     return cube_k, cube_residual, sk
44 end
```

Out[18]: svdize\_cube

Revert Submit

What does the movie represented by  $A_1$  look like? Let's find out! First we compute  $A_1$  using the function we just defined.

```
In [19]: 1 # Desired rank
2 k = 1
3 # Compute truncated SVD-version of movie
4 movie_cube_k, movie_cube_residual, _ = svdize_cube(movie_cube, k)
5
6 movie_cube_k = rescale_zero_one(movie_cube_k)
7 movie_cube_residual = rescale_zero_one(movie_cube_residual);
```

▶ Run

Done! Now let's examine the frames of the rank- $k$  approximation!

```
In [20]: 1 nframes = size(movie_cube, 3)
2 @manipulate for frame = 325
3     plotframe(
4         movie_cube_k[:, :, frame],
5         title = "Rank $k approx: Frame $frame of $nframes",
6         clim=(0, 1)
7     )
8 end
9 # Enter values in the box below to explore the video
```

▶ Run

Out[20]: frame 325

Rank 1 approx: Frame 325 of 650

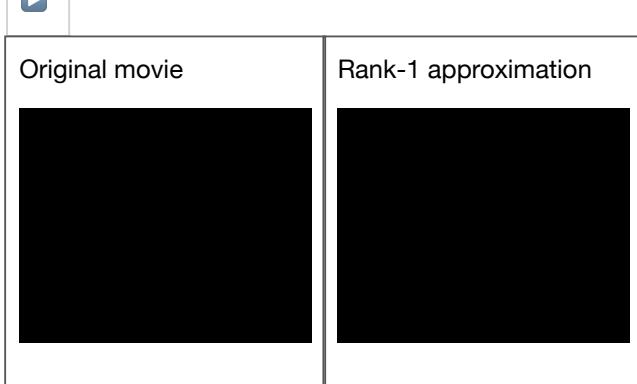


Perhaps you're seeing that something unexpected (it was to me, too!) has happened. Let's see the original movie and the movie of the rank- $k$  approximation side-by-side.

```
In [21]: 1 playvideo(  
2     [movie_cube, movie_cube_k],  
3     ["Original movie", "Rank-1 approximation"],  
4     frames_per_second=120  
5 )
```

Run

Out[21]:



What all happens in the rank-1 approximation movie? (select **all** that apply)



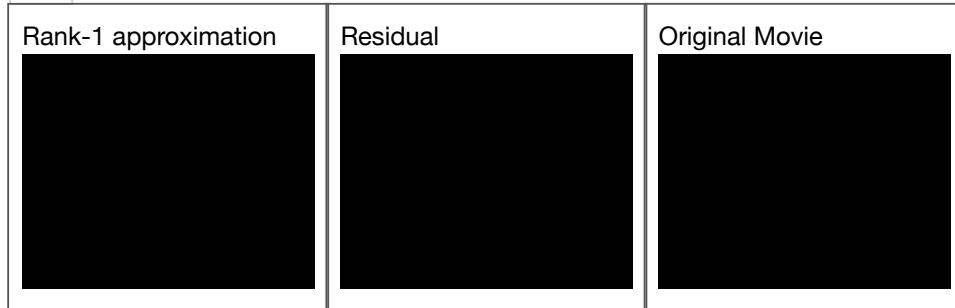
- There are people walking around, as in the original movie .
- Lights go on and off.
- The middle light in the background stays on throughout the scene.
- The four lights go from on to off, then come back on and go completely dark as before.
- The four lights go from on to off, then come back on, but faded (they do not go as dark as before).

Submit

Where did the people go? Let us examine the residual - after all, `movie_cube = movie_cube_k + movie_cube_residual` !

```
In [22]: 1 playvideo(
2     [movie_cube_k, movie_cube_residual, movie_cube],
3     ["Rank- $k$  approximation", "Residual", "Original Movie"],
4     frames_per_second=120
5 )
```

Out[22]: 325



Let's take a closer look at a fixed frame in each movie.

```
In [23]: 1 # An interesting frame to compare
2 frame = 250
3
4 hbox(
5     plotframe(
6         movie_cube[:, :, frame],
7         title = "Original movie: frame = $frame",
8         clim=(0, 1)
9     ),
10    plotframe(
11        movie_cube_k[:, :, frame],
12        title="Rank- $k$  approx: frame = $frame",
13        clim=(0, 1)
14    ),
15    plotframe(
16        movie_cube_residual[:, :, frame],
17        title="Rank- $k$  residual: frame = $frame",
18        clim=(0, 1)
19    ),
20)
```

Out[23]:



```
In [24]: 1 frame = 20
2
3 ▼ hbox(
4 ▼   plotframe(
5       movie_cube[:, :, frame],
6       title = "Original movie: frame = $frame",
7       clim=(0, 1)
8   ),
9 ▼   plotframe(
10      movie_cube_k[:, :, frame],
11      title="Rank-$k approx: frame = $frame",
12      clim=(0, 1)
13   ),
14 )
```

Run

Out[24]:



```
In [25]: 1 frame = 97
2
3 ▼ hbox(
4 ▼   plotframe(
5       movie_cube[:, :, frame],
6       title = "Original movie: frame = $frame",
7       clim=(0, 1)
8   ),
9 ▼   plotframe(
10      movie_cube_k[:, :, frame],
11      title="Rank-$k approx: frame = $frame",
12      clim=(0, 1)
13   ),
14 )
```

Run

Out[25]:



```
In [26]: 1 frame = 336
2
3 ▼ hbox(
4 ▼   plotframe(
5       movie_cube[:, :, frame],
6       title = "Original movie: frame = $frame",
7       clim=(0, 1)
8   ),
9 ▼   plotframe(
10      movie_cube_k[:, :, frame],
11      title="Rank-$k approx: frame = $frame",
12      clim=(0, 1)
13   ),
14 )
```

▶ Run

Out[26]:

Original movie: frame = 336



Rank-1 approx: frame = 336



```
In [27]: 1 frame = 650
          2
          3 ▾ hbox(
          4 ▾   plotframe(
          5       movie_cube[:, :, frame],
          6       title = "Original movie: frame = $frame",
          7       clim=(0, 1)
          8   ),
          9 ▾   plotframe(
          10      movie_cube_k[:, :, frame],
          11      title="Rank-$k approx: frame = $frame",
          12      clim=(0, 1)
          13   ),
          14 )
```

Run

Out[27]:



What does the rank-1 component capture?



- The "background" parts of the scene not containing people .
- The people.

Submit

## ▼ 3.2 What information do the right singular vectors represent?

We will soon discuss why the rank-1 component captures the background. It has to do with the fact that there is a low-rank matrix hiding inside the video, if we view things just right. Recall that the SVD returns the left and right singular vectors. So we may deduce the structure of the low-rank matrix hiding inside this movie.

Let us first try to understand what the right singular vectors of the decomposed movie convey. That will give us the important clue for unlocking this puzzle; once we do, we'll start seeing (induced) low-rank matrices (almost) everywhere!

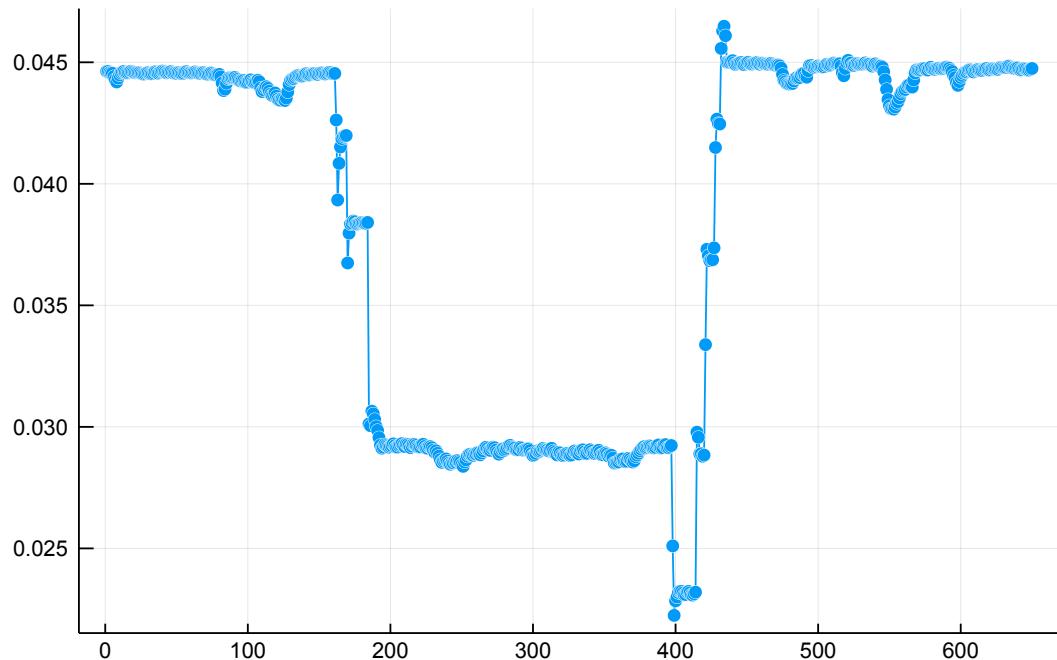
We begin by examining first right singular vector,  $V[:, 1]$ , more closely.

```
In [28]: 1 # Largest right singular vector
2 UsV = svds(reshape(movie_cube, :, size(movie_cube, 3)), nsv=1)[1]
3
4 # Generically, the sign of V is arbitrary, but it's easier to interpret
5 # our case) if the first element is positive.
6 positify(V) = sign(V[1, 1]) * V
7 V = positify(UsV.V)
8
9 plot(
10     V[:, 1],
11     linestyle=:solid,
12     marker=:circle,
13     title="Largest singular vector of movie matrix",
14     legend=:bottomleft
15 )
```

Run

Out[28]:

Largest singular vector of movie matrix



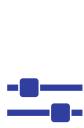
length(V[:, 1]) equals? (select all that apply)



- the number of frames in the movie.
- 650.
- 128 x 160 -- the total number of elements in each frame of the video.

Submit

length(U[:, 1]) equals? (select all that apply)



- 650.
- 128 x 160.

- the total number of elements in each frame (once it has been reshaped into a column vector).

 Submit

Do you notice something unusual in the behavior around elements 194 and 450? Recall that the  $A$  matrix has pixel indices on the rows and time on the columns. Thus,  $V$  encodes time behavior, and has dimensions equal to the number of frames in the movie. Let us scroll through the movie one frame at a time to see if we can identify the physical phenomenon that occurs in the movie when the elements of  $V[:, 1]$  "jump" (up or down).

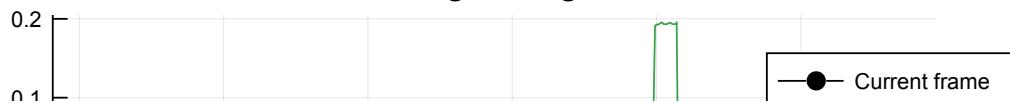
```
In [29]: 1 ▾ joinframes(frame::Integer, vids::Tuple) = hcat((vid[:, :, frame] for
2
3     @manipulate for frame = (100, 194, 397, 450)
4         # Plot first right singular vector
5         p1 = plot(
6             V,
7             linestyle=:solid,
8             title="First right singular vector"
9         )
10
11     # Denote current frame
12     plot!(
13         (frame, V[frame, 1]),
14         marker=:circle, # current point
15         color=:black,
16         label="Current frame"
17     )
18
19     # Plot movie frames
20     p2 = plotframe(
21         joinframes(
22             frame,
23             (movie_cube,
24             movie_cube_k,
25             movie_cube_residual)
26         ),
27         title="Movie | Rank-k approx. | Residual",
28         clim=(0, 1)
29     )
30     vbox(plot(p1, p2; layout=Plots.grid(2, 1, heights=[0.5, 0.5])))
31
32 end
33
34 # Change frame to view different frames
```

▶ Run

Out[29]: frame

100 194 397 450

## First right singular vector



Frames 194 and 450 correspond (roughly) to the times when (select **all** that apply)



- The four lights go from off to on.
- The middle light in the background stays on.
- The four lights remain on.
- Nothing changes.

Submit

The right singular vector is encoding the "on-off-on" behavior of the front row of lights! Note that all elements of  $V[:, 1]$  are positive, so the "off" portion corresponds to 0 in the appropriate segment of the vector.

Let us now compare the low-rank decompositions obtained by setting  $k = 1$  versus  $k = 2$ , as in the next cell. When  $k = 2$ , we will also plot  $V[:, 2]$  and try to understand what it captures.

```
In [30]: 1 # one-time computation
2
3 nframes = size(movie_cube, 3)
4
5 back_1, res_1, _ = svdize_cube(movie_cube, 1)
6 back_2, res_2, _ = svdize_cube(movie_cube, 2)
7
8 V_1 = positify(svds(reshape(movie_cube, :, nframes); nsv=1)[1].V)
9 V_2 = positify(svds(reshape(movie_cube, :, nframes); nsv=2)[1].V);
```

Run

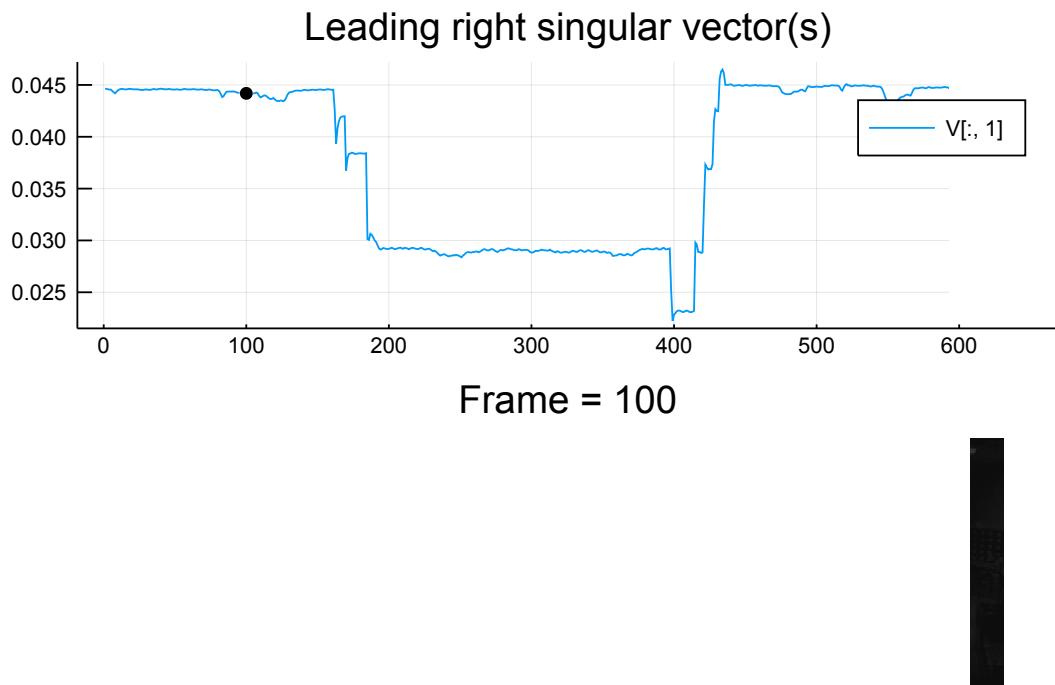
```
In [31]: 1 # interactive display
2 @manipulate for k = (1, 2), frame = (100, 194, 397, 450)
3     if k == 1
4         back_k = back_1
5         res_k = res_1
6         V_k = V_1
7     elseif k == 2
8         back_k = back_2
9         res_k = res_2
10        V_k = V_2
11    end
12
13    # Plot right singular vector(s)
14    p1 = plot(
15        V_k;
16        linestyle=:solid,
17        title="Leading right singular vector(s)",
18        label=["V[:, 1]", "V[:, 2]"]
19    )
20
21    # Denote current frame
22    for kk in 1:k
23        plot!(
24            (frame, V_k[frame, kk]);
25            marker=:circle, # current point
26            color=:black,
27        )
28    end
29
30    # Plot movie frames
31    p2 = plotframe(
32        joinframes(
33            frame,
34            (movie_cube,
35            back_k,
36            res_k)
37        ),
38        title="Frame = $frame",
39        clim=(0, 1)
40    )
41
42    plot(
43        p1, p2;
44        layout=Plots.grid(2, 1; heights=[0.5, 0.5])
45    )
46 end
47
48 # Change k to use a different rank approximation
49 # Change frame to view different frames
```

Out[31]: k

1 2

frame

100 194 397 450



Now let's plot the individual frames of the decomposition for  $k = 2$ .

```
In [32]: 1 # Choose rank
2 k = 2
3
4 # Decompose movie
5 movie_cube_k, movie_cube_residual, _ = svdsize_cube(movie_cube, k);
6 UsV = svds(reshape(movie_cube, :, size(movie_cube, 3)), nsv=k)[1]
7 V = positivify(UsV.V);
```

Run

```
In [33]: 1 # An interesting frame
2 frame = 250
3
4 hbox(
5     # Original movie
6     plotframe(
7         movie_cube[:, :, frame],
8         title="Original movie: frame $frame",
9         clim=(0, 1)
10    ),
11    # Rank-k approx
12    plotframe(
13        movie_cube_k[:, :, frame],
14        title="Rank $k approx: frame $frame",
15        clim=(0, 1)
16    ),
17    # Residual
18    plotframe(
19        movie_cube_residual[:, :, frame],
20        title="Rank $k residual: frame $frame"
21    )
22 )
```

**Run**

Out[33]:



Compared to  $k = 1$ , the  $k = 2$  low-rank decomposition (select **all** that apply)



- better captures the on-off-on behavior of the row of four lights.
- captures more accurately the portion of the movie where only one light is on.
- has a residual (with the people moving around) with a higher image contrast.

**Submit**

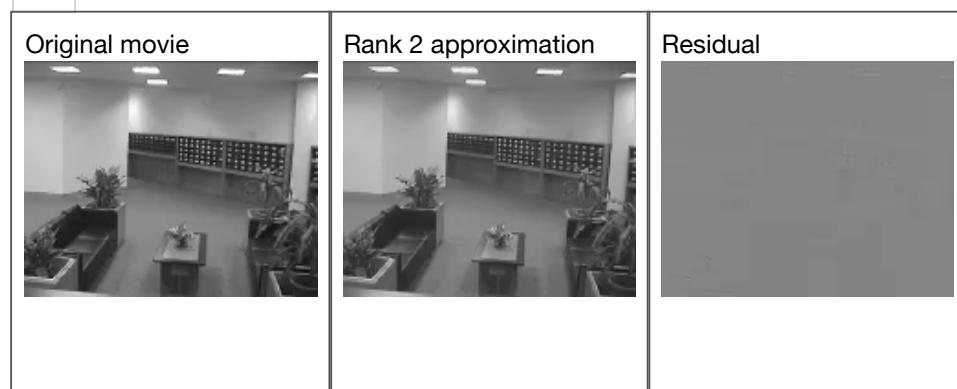
Setting  $k = 2$  does indeed better capture the background than  $k = 1$ . Particularly we are better able to capture the on-off-on behavior of the row of lights in the front. As before, the people have been removed in the low rank approximation and instead appear in the residual matrix. The movie in the next cell shows that this is true for all the frames.

```
In [34]: 1  println("Original movie vs Rank $k approximation vs Residual")
2  playvideo(
3  [movie_cube, rescale_zero_one(movie_cube_k), rescale_zero_one(mov
4  ["Original movie", "Rank $k approximation", "Residual"], frames_p
5 )
```

▶ Run

Original movie vs Rank 2 approximation vs Residual

Out[34]:

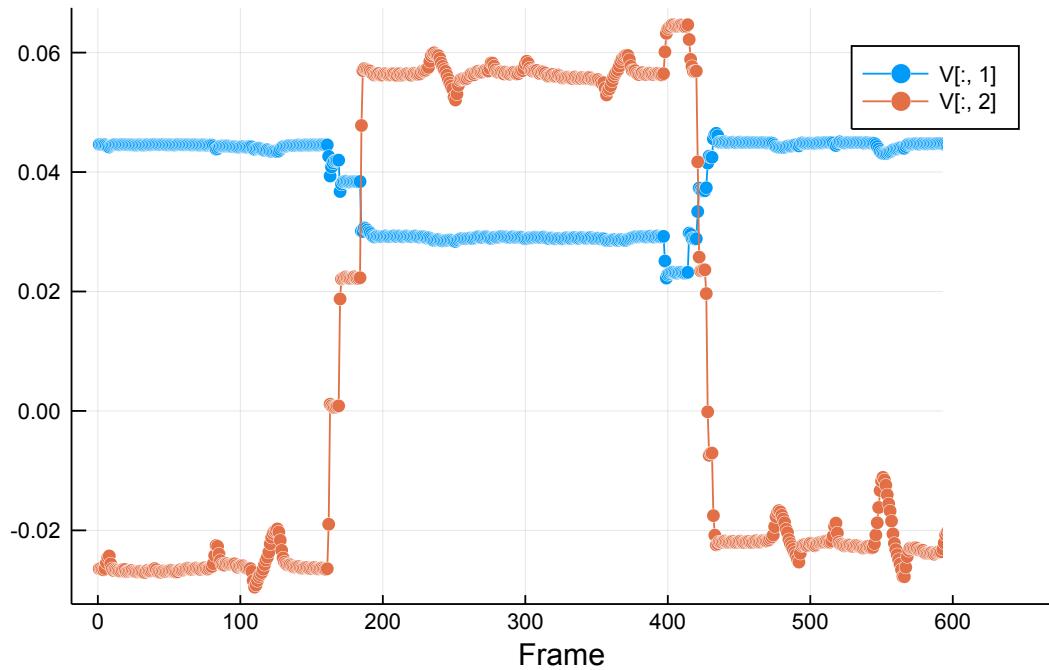


Let us look at the top two right singular vectors using the command in the next cell.

```
In [35]: 1 plot(
2     V,
3     label=[ "V[:, $i]" for i in collect(1:k) ],
4     title="$k leading right singular vectors of movie matrix",
5     xaxis="Frame",
6     marker=:circle,
7 )
```

Out[35]:

2 leading right singular vectors of movie matrix



Notice the step-function-like behavior in the singular vectors! The elements of both the first and second singular vectors jump up (or down) at around index 195 and back (or up) at around index 420.

These time instances correspond to the portions of the movie where the front lights lights go from on to off and then from off to on. We verify this observation using by slicing through the video around these frames as in the next cell.

```
In [36]: 1 @manipulate for frame = (150, 196, 300, 420, 500)
2     plotframe(
3         movie_cube[:, :, frame],
4         title = "Original movie: frame = $frame",
5     )
6 end
7
8 # Change frame to view different frames
```

Out[36]: frame

150 196 300 420 500

Original movie: frame = 150



What would `movie_cube_k` capture if we set `k = r`, where `r = rank(movie_cube)`? (select all that apply)

- Increasing `r` decreases the error, so we will better separate the background from the foreground (people moving).



movie\_cube\_k = movie\_cube by the definition of the SVD, so movie\_cube\_residual = 0 , and we would not be doing any background subtraction

Increasing  $r$  monotonically decreases the error between movie\_cube\_k and movie\_cube , but what we are really interested in is the ~~error between~~ movie\_cube\_k and the background component of the scene.

### 3.3 The low-rank matrix hiding inside the video

There are two main things happening in the video that can be captured via a low-rank matrix.

Let  $X_1$  represent the  $128 \times 160$  matrix (or image) corresponding to what the room looks like when only the middle light in the background is on, and the four lights in the foreground are off. Let  $x_1$  denote the vector formed by reshaping  $X_1$  into a  $128 \times 160$  dimensional vector.

Let  $X_2$  represent the  $128 \times 160$  matrix corresponding to what the room looks like when only the four lights in the foreground are on. Let  $x_2$  denote the vector formed by reshaping  $X_2$  into a  $128 \times 160$  dimensional vector.

See below for what  $X_1$  (left) and  $X_2$  (right) look like.



X1



X2

Let  $B_t$  denote the  $128 \times 160$  matrix corresponding to the background at time instance  $t$ . Then assuming a linear illumination model, we can express  $B_t$  as

$$B_t = y_{1t}X_1 + y_{2t}X_2, \quad (1)$$

where  $y_{1t}$  denotes the relative intensity of the middle light in the background, and  $y_{2t}$  denotes the relative intensity of the foreground row of four lights. If  $y_{1t} = 0$ , then only the lights in the foreground are on; if  $y_{2t} = 0$ , then only the light in the background is on.

Let us see what various linear combinations of  $X_1$  and  $X_2$  look like. We'll vary  $y_1$  and  $y_2$  and see what happens.

```
In [37]: 1 frontback = load("front_back.jld")
2 ▼ front = frontback["front"]
3 ▼ back = frontback["back"]
4
5 @manipulate for y1 in (1.0, 0.5, 0.0), y2 in (0.0, 1.0)
6     plotframe(y1 * back + y2 * front)
7 end
```

Out[37]:  
y1  
y2

|     |     |     |
|-----|-----|-----|
| 1.0 | 0.5 | 0.0 |
| 0.0 | 1.0 |     |



The background consists of a scene where the middle light is always on, and the four lights towards the front go on and off. Which of the following statements accurately connects matrix language to what we can see with our eyes? (select all apply)



- $y_1 = 0, y_2 = 0$  means only the middle light towards the back is on.
- $y_1 = 1, y_2 = 0$  means only the middle light towards the back is on.
- $y_1 = 1, y_2 = 1$  means only the front four lights are on.
- $y_1 = 0, y_2 = 1$  means only the front four lights are on.
- $y_1 = 1, y_2 = 1$  corresponds to all the lights being on.

Let  $b_t$  denote the vector formed by reshaping  $B_t$  into a  $128 \times 160$  vector. Then from Eq. (???) we have that

$$b_t = y_{1t}x_1 + y_{2t}x_2 = \begin{bmatrix} x_1 & x_2 \end{bmatrix} \begin{bmatrix} y_{1t} \\ y_{2t} \end{bmatrix}.$$

Let  $B$  denote the  $mn \times T$  matrix formed by stacking  $b_1, \dots, b_T$  alongside each other. Then

$$B = \begin{bmatrix} x_1 & x_2 \end{bmatrix} \begin{bmatrix} y_{11} & \dots & y_{1T} \\ y_{21} & \dots & y_{2T} \end{bmatrix} = x_1 y_1^T + x_2 y_2^T,$$

where  $y_1$  and  $y_2$  are  $T$  dimensional vectors defined as

$$y_1 = [y_{11} \dots y_{1T}]^T,$$

and

$$y_2 = [y_{21} \dots y_{2T}]^T,$$

The final step in uncovering this rank-2 background matrix hidden inside the video is assigning (relative) values to the elements of  $y_1$  and  $y_2$  so we can capture the dynamics of how the lights switch on or off.

Suppose  $x_1$  and  $x_2$  denote the components of the scene described above. Then a plausible set of values for  $y_1$  and  $y_2$  (that captures the dynamics but not actual intensities of lights) is



$y_1 = [1, 1, \dots, 1]^T, y_2 = [0, 0, \dots, 1, \dots, 1, 0, \dots, 0]^T$

$y_1 = [1, 1, \dots, 1]^T, y_2 = [1, 1, \dots, 0, \dots, 0, 1, \dots, 1]^T$

$y_1 = [1, 1, \dots, 0, \dots, 0, 1, \dots, 1]^T, y_2 = [1, 1, \dots, 1]^T$

Submit

Indeed! If  $y_1$  is all ones, the middle light will be on all the time, which matches what we observe. And if  $y_2$  is a series of ones, followed by zeros, followed by ones again, then the front four lights will go from on to off and then back on.

We could vary the elements in  $y_1$  and  $y_2$  with time (instead of setting them to either 1 or 0) to encode variation in the intensity of the light; this would not change the rank of the resulting background matrix.

The presence of this underlying latent (or hidden) rank-2 matrix inside the video is why the rank-two approximation better captures the dynamics of the background. Note that the SVD has captured a dynamic, *time-varying* background. We could have computed a mean matrix by averaging all the images in the video. That would remove the people, but it would be static and would **not** capture the behavior of the lights!

Spotting this low-rank behavior allowed us to use the SVD to powerful effect.

What do we not yet fully understand? (select **all** that apply)

- If the people are in the residual matrix, and we can think of the movie matrix as background + people, then why are the singular vectors of the movie matrix so close to the singular vectors of the people-free background matrix?



- Can we determine the rank of the background matrix automatically, without going through the reasoning process we just did?
  - Would we have recovered the background matrix if noise was added to the movie?
  - Would we have recovered the background matrix if we randomly zeroed out (or deleted) pixels in the movie?
  - If we continually add noise and delete entries, at what point will we become unable to accurately recover the low-rank background matrix?
  - Everything is illuminated -- there are no questions left to ask!
- The  $y_1$  and  $y_2$  we modeled are linearly independent but not orthogonal, whereas the right singular vectors from the SVD are orthogonal. Is there a technique for expressing the movie matrix as a sum of outer products with non-orthogonal components that might model  $y_1$  and  $y_2$  better?
- I am still confused, but at a higher level, about more important questions!

Submit

We will soon discuss insights from random matrix theory and matrix perturbation theory that address all of the above questions! We will also discuss other algorithms for decomposing (or factorizing) a matrix.

## 4 Summary

We saw how the SVD allows us to capture low-rank matrices. The ability to spot low rank matrices is a critical skill in computational data science. Rearrangement of the data can induce low rank structure, which the SVD can then expose.

## 5 Additional Exercises

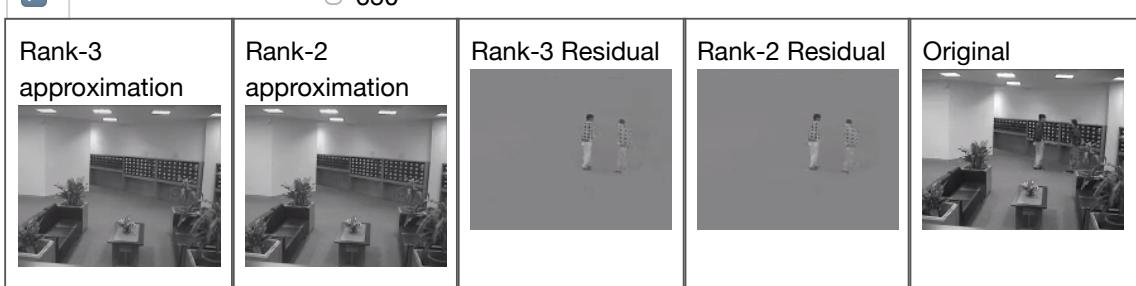
### 5.1 Comparing the $k = 2$ vs $k = 3$ decompositions

Perform a truncated low rank decomposition of the video, as above, except with  $k = 3$ . Display the video of the original and residual next to each other, and compare it to the setting when  $k = 2$ .

```
In [74]: 1 ## TODO: Compare videos of the k = 2 vs k = 3 reduced-rank videos.  
2  
3 movie_cube_2, movie_cube_residual_2, _ = svdize_cube(movie_cube, 2)  
4 movie_cube_2 = rescale_zero_one(movie_cube_2)  
5 movie_cube_residual_2 = rescale_zero_one(movie_cube_residual_2);  
6  
7 movie_cube_3, movie_cube_residual_3, _ = svdize_cube(movie_cube, 3)  
8 movie_cube_3 = rescale_zero_one(movie_cube_3)  
9 movie_cube_residual_3 = rescale_zero_one(movie_cube_residual_3);  
10  
11 playvideo(  
12     [movie_cube_3, movie_cube_2, movie_cube_residual_3, movie_cube_re  
13     ["Rank-3 approximation", "Rank-2 approximation", "Rank-3 Residual"  
14     frames_per_second=120  
15 )
```

▶ Run

Out[74]:



```
In [39]: 1 ## TODO: Compare representative frames of the k = 2 vs k = 3
2 # reduced-rank videos. Hint: use @manipulate
3 @manipulate for frame = (150, 196, 300, 420, 500)
4 hbox(
5     # Original movie
6     plotframe(
7         movie_cube_2[:, :, frame],
8         title="Rank 2 approx: frame $frame",
9         clim=(0, 1)
10    ),
11    # Rank-k approx
12    plotframe(
13        movie_cube_3[:, :, frame],
14        title="Rank 3 approx: frame $frame",
15        clim=(0, 1)
16    ),
17 )
18 end
```

Run

Out[39]: frame

150 196 300 420 500

Rank 3 approx: frame 150



Which better captures the background? Display some representative frames.

```
In [40]: 1 frame = 171
2 ▼ hbox(
3     # Original movie
4 ▼ plotframe(
5         movie_cube_2[:, :, frame],
6         title="Rank 2 approx: frame $frame",
7         clim=(0, 1)
8     ),
9     # Rank-k approx
10 ▼ plotframe(
11         movie_cube_3[:, :, frame],
12         title="Rank 3 approx: frame $frame",
13         clim=(0, 1)
14     ),
15 )
16
```

▶ Run

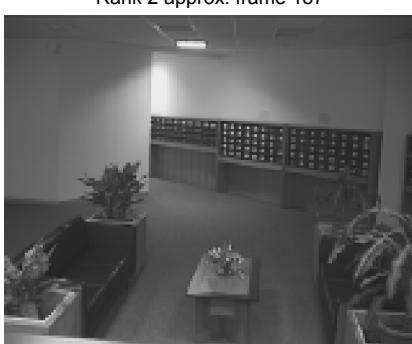
Out[40]:



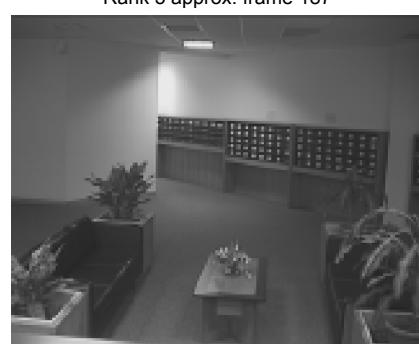
```
In [41]: 1 frame = 187
2 ▼ hbox(
3     # Original movie
4 ▼ plotframe(
5         movie_cube_2[:, :, frame],
6         title="Rank 2 approx: frame $frame",
7         clim=(0, 1)
8     ),
9     # Rank-k approx
10 ▼ plotframe(
11         movie_cube_3[:, :, frame],
12         title="Rank 3 approx: frame $frame",
13         clim=(0, 1)
14     ),
15 )
```

▶ Run

Out[41]:



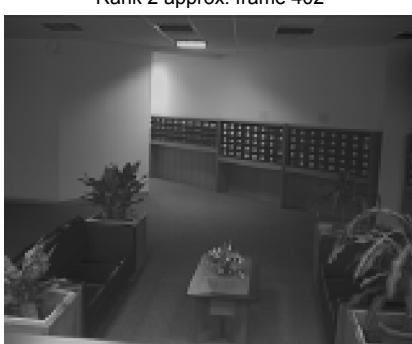
Rank 3 approx: frame 187



```
In [42]: 1 frame = 402
2 ▼ hbox(
3     # Original movie
4 ▼ plotframe(
5         movie_cube_2[:, :, frame],
6         title="Rank 2 approx: frame $frame",
7         clim=(0, 1)
8     ),
9     # Rank-k approx
10 ▼ plotframe(
11         movie_cube_3[:, :, frame],
12         title="Rank 3 approx: frame $frame",
13         clim=(0, 1)
14     ),
15 )
```

▶ Run

Out[42]:



```
In [43]: 1 frame = 425
2 ▼ hbox(
3     # Original movie
4 ▼ plotframe(
5         movie_cube_2[:, :, frame],
6         title="Rank 2 approx: frame $frame",
7         clim=(0, 1)
8     ),
9     # Rank-k approx
10 ▼ plotframe(
11         movie_cube_3[:, :, frame],
12         title="Rank 3 approx: frame $frame",
13         clim=(0, 1)
14     ),
15 )
```

Run

Out[43]:

Rank 2 approx: frame 425



Rank 3 approx: frame 425



Rank 3 approx better captures the background, which is clear if you see the residuals at the above frames. Rank 2 approx deletes some of the objects in the background like plants while rank 3 captures these better.

Edit / Render

## ▼ 5.2 Information contained in the third singular vector

Plot the first three right singular vectors. Remark on the characteristics of the third singular vector.

- Is it flat?
- Does it have "jumps"?
- Are the jumps smooth, or sharp?
- What is happening in the video when the jumps occur?

It is not flat, the singular vector have jumps.



The jumps are sharp.



The lights in the background turn on or off when the jumps occur.

Edit / Render

In [49]:

```

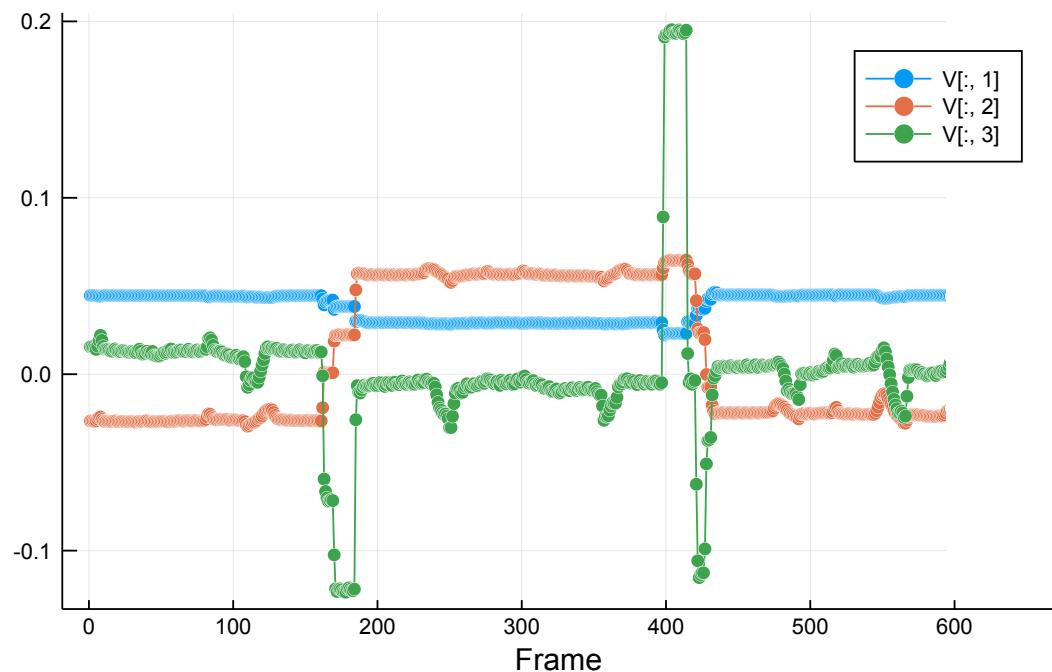
1 ##TODO: Plot the first three right singular vectors
2 UsV = svds(reshape(movie_cube, :, size(movie_cube, 3)), nsv=3)[1]
3 V = positivify(UsV.V);
4 plot(
5     V,
6     label=["V[:, $i]" for i in collect(1:3)],
7     title="$k leading right singular vectors of movie matrix",
8     xaxis="Frame",
9     marker=:circle,
10)
11

```

Run

Out[49]:

2 leading right singular vectors of movie matrix



```
In [72]: 1 ## TODO: display what happens in the video at time instances
2 # corresponding to jumps or anything interesting happening
3 # in the third right singular vector.
4 frame = 397
5 hbox(
6     # Original movie
7     plotframe(
8         movie_cube[:, :, frame],
9         title="Original",
10        clim=(0, 1)
11    ),
12    # Rank-k approx
13    plotframe(
14        movie_cube_3[:, :, frame],
15        title="Rank 3 approx: frame $frame",
16        clim=(0, 1)
17    ),
18 )
19
```



Out[72]:

Original



Rank 3 approx: frame 397



In the discussion earlier, we spoke about how the scene could be represented as  $x_1 y_1^T + x_2 y_2^T$ . If the scene is better described by a rank-3 matrix, that suggests that the light on/off phenomenon in the video (absent people) is actually better modeled by  $x_1 y_1^T + x_2 y_2^T + x_3 y_3^T$ .

What do  $x_3$  and  $y_3$  physically represent, if  $x_1, x_2$  and  $y_1, y_2$  are as described earlier?

Your answer here!





## ▼ 5.3 The singular value spectrum of the reshaped video matrix

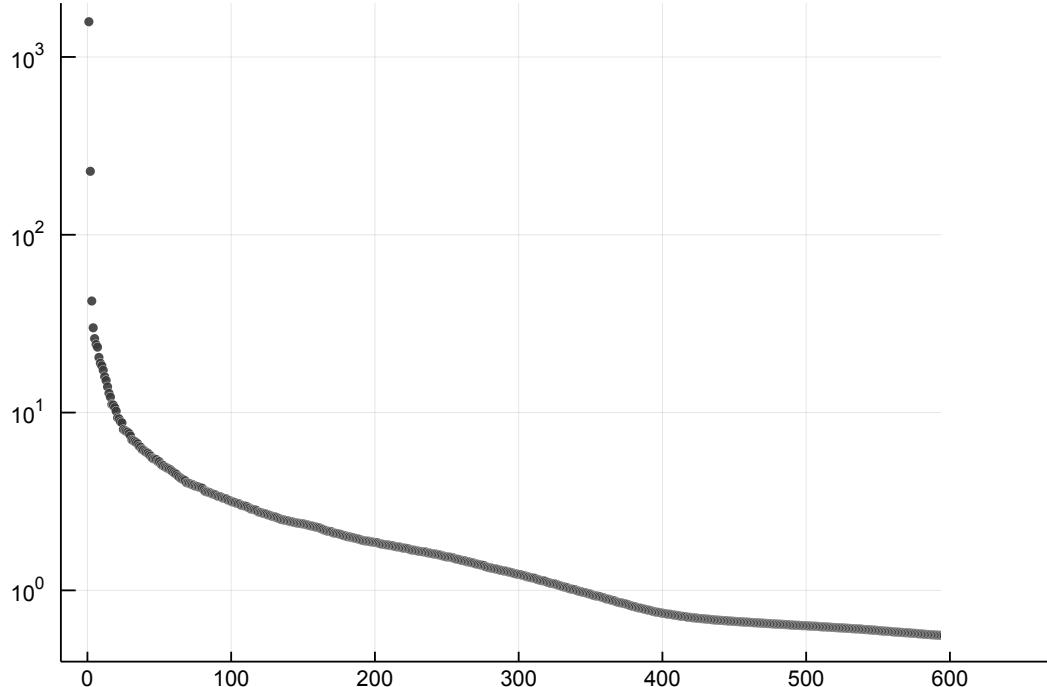
It turns out that the singular value spectrum of the video matrix contains information that could

have allowed us to determine the implicit rank ( $k = 3$ ) of the "light on/off" matrix. Plot the singular value spectrum in the cell below. Notice that many of the values appear grouped together, and a few appear to be well-separated from this group.

```
In [69]: 1 ##TODO: Plot the singular values of the reshaped video matrix.
2 # Tip: Use the `scatter` and `svdvals` commands. Consider plotting on
3 nframes = size(movie_cube, 3)
4 movie_matrix = reshape(movie_cube, :, nframes)
5
6 using LinearAlgebra: svdvals
7 s = svdvals(movie_matrix)
8
9 scatter(1:length(s), s;yscale=:log10, alpha=0.7, markersize=3, c=:bl
```

 Run

Out[69]:



It can be shown mathematically (and we will do so later) that one can estimate the rank of low-rank structure by counting the number of singular values that are relatively separate from the "continuous-looking portion" of the singular value spectrum.

Do you agree with this claim?

Yes



 Edit / Render

## 5.4 Comparing to the "mean frame"

What if instead of using the SVD, we just took the average of all the frames? How would that differ from our low-rank approximation?

```
In [65]: 1 # write code to compute the "mean frame movie" and
2 # display it alongside the original movie and a
3 # rank-k approximation
4 m,n,l = size(movie_cube)
5 movie_cube_mean = zeros(m,n,l)
6 movie_cube_mean[:, :, 1] = sum(movie_cube, dims = 3)/l
7 for i in 1:l
8     movie_cube_mean[:, :, i] = movie_cube_mean[:, :, 1]
9 end
10 playvideo(
11     [movie_cube, movie_cube_2, movie_cube_mean],
12     ["Original", "Rank-2 approximation", "Mean Movie"],
13     frames_per_second=120
14 )
```

Run

Out[65]:



```
In [67]: 1 plotframe(  
2         movie_cube_mean[:, :, 1],  
3         title="Mean frame",  
4         clim=(0, 1)  
5     )
```

Run

Out[67]:

Mean frame

